

NBA Daily Fantasy Basketball Predictor 🏀

A comprehensive machine learning pipeline for NBA daily fantasy sports (DFS) that implements advanced prediction models and optimization algorithms to generate profitable lineups. This project represents cutting-edge research in applying temporal machine learning techniques to DFS competitions.

🎯 Overview

This project addresses a fundamental research question: **Can advanced ML techniques that capture temporal dependencies in NBA data, combined with sophisticated optimization algorithms, generate consistent profits in DFS tournaments?**

Key innovations include:

- **Temporal Modeling:** Rolling-window training with LSTM/RNN and XGBoost
- **Advanced Feature Engineering:** 900+ features including DvP metrics, cross-season stats, and volatility indicators
- **Multi-Algorithm Optimization:** Genetic algorithms, integer linear programming, and constraint solvers
- **Comprehensive Backtesting:** Evaluation across 8 NBA seasons with real contest data
- **Salary-Based Stratification:** Targeted modeling for high-value players

🚀 Quick Start

Prerequisites

- Python 3.8+
- For Apple Silicon GPU acceleration: PyTorch with MPS support

Installation

```
bash

git clone https://github.com/tzachpach/final_project_dfs.git
cd final_project_dfs
pip install -r requirements.txt
```

Running the Pipeline

```
bash
```

```
python src/experiments/run_experiments.py
```

This command initiates the comprehensive experimental pipeline that:

1. Runs multiple ML models (XGBoost, LSTM/RNN) with various configurations
2. Tests different feature engineering strategies
3. Evaluates multiple optimization algorithms
4. Generates performance metrics and comparisons
5. Outputs results for analysis

Workflow Architecture

```
mermaid
```

```
graph TD
```

```
A[Raw NBA Data] --> B[run_experiments.py]
B --> C[Experimental Pipeline]
C --> D[Model Training & Evaluation]
D --> E[XGBoost Models]
D --> F[LSTM/RNN Models]
E --> G[Fantasy Point Predictions]
F --> G
G --> H[Optimization Algorithms]
H --> I[Genetic Algorithm]
H --> J[Integer Linear Programming]
H --> K[PuLP Solver]
I --> L[Optimal DFS Lineup]
J --> L
K --> L
L --> M[Contest Backtesting]
M --> N[Profit/Performance Analysis]
```

Core Components

1. Experimental Framework (`src/experiments/run_experiments.py`)

- **Main entry point** for the entire pipeline
- Orchestrates systematic testing of different model configurations
- Manages hyperparameter grid searches

- Coordinates model training, evaluation, and comparison

2. Data Preprocessing (`src.preprocessing`)

- `merge_all_seasons()`: Combines multiple season CSV files
- `preprocess_all_seasons_data()`: Standardizes columns and cleans data
- Handles missing values and inconsistent formatting

3. Feature Engineering (`src.data_enrichment`)

- `add_time_dependent_features_v2()`:
 - Rolling averages (3, 5, 10 game windows)
 - Performance differentials
 - Trend indicators
 - Volatility metrics
- `add_running_season_stats()`: Cumulative season statistics
- `add_last_season_data_with_extras()`: Cross-season continuity features
- **Defense vs Position (DvP)** metrics for matchup analysis

4. Prediction Models

XGBoost

- **Module:** `src.predict_fp_xgb_daily`
- **Features:**
 - Daily and weekly rolling predictions
 - Salary quantile binning (top 10%, 25%, etc.)
 - Multi-target vs single-target modes
 - Feature reduction strategies (SelectKBest, PCA)

LSTM/RNN

- **Module:** `src.predict_fp_rnn_*`
- **Features:**
 - Sequential modeling with lookback windows
 - Per-player scaling
 - Hyperparameter tuning for architecture

- Multi-target prediction capability

5. Lineup Optimization

- **Genetic Algorithm** (`src.lineup_genetic_optimizer`)
 - Population-based optimization
 - Custom mutation/crossover strategies
 - Constraint-aware evolution
- **Integer Linear Programming** (ILP)
 - Google OR-Tools implementation
 - Guaranteed optimal solutions
 - Exact constraint modeling
- **PuLP Solver**
 - Alternative ILP implementation
 - Enhanced debugging capabilities



Contest Outcome Analysis

The pipeline includes comprehensive backtesting against real DFS contest data to evaluate profitability:

Setup

1. Ensure you have contest data: `fanduel_nba_contests.csv`
2. Generated lineups from the experimental pipeline
3. Run analysis scripts or notebooks for performance evaluation

Key Metrics

- **Win Rate:** Percentage of contests where lineup would've won
- **Cash Rate:** Percentage of contests where lineup would've cashed (finished above payout threshold)
- **ROI:** Return on investment over time
- **Profit Analysis:** Daily and cumulative profit tracking

Performance Evaluation

The experimental framework automatically:

- Compares predicted vs actual fantasy points
- Tests lineup performance across thousands of historical contests
- Evaluates different model configurations for profitability
- Identifies optimal strategies for consistent returns

Sample Results (from paper)

- Tested across 8 NBA seasons (2017-2024)
- 17,811 contests per season average
- \$234M total contest participation per season
- Multiple model architectures compared for optimal performance

Project Structure

```

final_project_dfs/
├── src/
│   ├── experiments/
│   │   └── run_experiments.py    # Main entry point – orchestrates all experiments
│   ├── preprocessing/          # Data cleaning and merging
│   ├── data_enrichment/        # Feature engineering modules
│   ├── predict_fp_xgb_daily/    # XGBoost prediction module
│   ├── predict_fp_rnn_*/       # RNN/LSTM prediction modules
│   └── lineup_genetic_optimizer/ # Optimization algorithms (GA, ILP, PuLP)
├── data/                       # NBA game logs and DFS contest data
├── output_csv/                 # Final lineup predictions and results
├── notebooks/
│   └── analyze_contests_vs_pred.ipynb # Contest analysis
├── models/                     # Saved model files
├── requirements.txt            # Python dependencies
└── main.py                     # Legacy single-run script (deprecated)

```

Configuration

Experimental Pipeline Configuration

The `run_experiments.py` script orchestrates extensive hyperparameter searches across:

Model Parameters

- **XGBoost:** Tree depth, learning rate, estimators, feature reduction methods
- **LSTM/RNN:** Hidden layers, dropout, batch size, learning rate, lookback windows

Training Strategies

- **Rolling Windows:** Daily vs Weekly training periods
- **Salary Binning:** Top 10%, 25%, or full player pool
- **Prediction Modes:** Multi-target (individual stats) vs Single-target (direct FP)
- **Feature Reduction:** SelectKBest, PCA, or no reduction

Optimization Algorithms

The pipeline tests all three optimization methods:

- Genetic Algorithm with configurable population size and generations
- Integer Linear Programming with exact solutions
- PuLP solver for comparison and validation

Customization

To modify experimental parameters, edit the configuration in `run_experiments.py`:

```
python
```

```
# Example: Adjust model hyperparameters
```

```
xgb_params = {  
    'max_depth': [3, 5, 7],  
    'eta': [0.01, 0.05, 0.1],  
    'n_estimators': [100, 200, 300]  
}
```

```
# Example: Change optimization settings
```

```
ga_config = {  
    'population_size': 100,  
    'generations': 50,  
    'mutation_rate': 0.1  
}
```



Feature Details

Temporal Features

- **Rolling Statistics:** 3, 5, and 10-game windows for all major stats
- **Volatility Metrics:** Standard deviation of recent performances
- **Trend Indicators:** Performance direction over last N games

- **Rest Days:** Games missed, back-to-back indicators
- **Lagged Features:** Day-to-day performance differences

Matchup-Specific Features

- **Defense vs Position (DvP):** How opposing teams defend specific positions
- **Pace Metrics:** Game tempo expectations
- **Opponent Defensive Ratings:** Team-specific defensive efficiency
- **Historical H2H Performance:** Player performance vs specific teams

Cross-Season Features

- Previous season averages and totals
- Career progression metrics
- Double-double and triple-double frequencies
- Exceptional game counts (20+ points, 10+ assists, etc.)

Advanced Metrics

- **Usage Rate:** Percentage of team plays while on court
- **Efficiency Ratings:** Offensive/Defensive/Net ratings
- **True Shooting Percentage:** Shooting efficiency accounting for 3PT and FT
- **Player Impact Estimate (PIE):** Overall contribution metric

Salary-Based Stratification

The experimental framework automatically bins players by salary percentiles:

- **Elite Tier:** Top 10% by salary (stars and superstars)
- **High Tier:** 75-90th percentile (solid starters)
- **Mid Tier:** 50-75th percentile (role players)
- **Value Tier:** Below 50th percentile (bench players)

This stratification ensures models focus on DFS-relevant players and capture tier-specific patterns.

Research Findings & Roadmap

Key Research Questions (from paper)

1. Can advanced ML techniques generate consistent profits in NBA DFS contests?

2. **How do temporal dependencies and feature engineering impact prediction accuracy?**
3. **Which optimization algorithms perform best under real-world constraints?**

Experimental Results Summary

The comprehensive experimental pipeline evaluates:

- **Model Architectures:** XGBoost vs LSTM/RNN performance
- **Feature Engineering:** Impact of rolling windows, DvP metrics, and salary binning
- **Optimization Methods:** Genetic Algorithm vs ILP solver effectiveness
- **Temporal Strategies:** Daily vs weekly prediction windows
- **Financial Performance:** Win rates, cash rates, and ROI across thousands of contests

Phase 1: Core Implementation

- ☒ Comprehensive experimental pipeline
- ☒ XGBoost and RNN implementations
- ☒ Multiple optimization algorithms
- ☒ Historical contest backtesting
- ☒ Advanced feature engineering

Phase 2: Advanced Research

- ☐ Real-time injury and news integration
- ☐ Multi-lineup generation for GPP contests
- ☐ Ownership projection models
- ☐ Bayesian hyperparameter optimization
- ☐ Deep reinforcement learning for lineup selection

Phase 3: Production Deployment

- ☐ Cloud-based model training infrastructure
- ☐ Real-time API for live contests
- ☐ Automated bankroll management
- ☐ Mobile companion application
- ☐ Integration with multiple DFS platforms

Contributing

We welcome contributions! Please follow these steps:

1. Fork the repository

2. Create a feature branch (`git checkout -b feature/AmazingFeature`)
3. Commit your changes (`git commit -m 'Add some AmazingFeature'`)
4. Push to the branch (`git push origin feature/AmazingFeature`)
5. Open a Pull Request

Code Style

- Follow PEP 8 guidelines
- Add docstrings to all functions
- Include unit tests for new features



License

This project is licensed under the MIT License - see the [LICENSE](#) file for details.



Acknowledgments

- NBA data sources and APIs
- DFS community for strategy insights
- Open source ML libraries (XGBoost, PyTorch, scikit-learn)



Contact

Tzach Pach - [@tzachpach](#)

Project Link: https://github.com/tzachpach/final_project_dfs

Disclaimer: This tool is for educational and research purposes. Always gamble responsibly.